

Importing necessary libraries -

```
1 import pandas as pd
2 import numpy as np
3 import seaborn as sns
4 from scipy import stats
5 import matplotlib.pyplot as plt
6 from sklearn.linear_model import LogisticRegression
7 from sklearn import metrics
8 from sklearn.metrics import confusion_matrix
9 from sklearn.metrics import classification_report
10 from sklearn.metrics import roc_auc_score
11 from sklearn.metrics import roc_curve
12 from sklearn.metrics import precision_recall_curve
13 from sklearn.model_selection import train_test_split, KFold, cross_val_score
14 from sklearn.preprocessing import MinMaxScaler
15 from sklearn.metrics import (
16     accuracy_score, confusion_matrix, classification_report,
17     roc_auc_score, roc_curve, auc,
18     plot_confusion_matrix, plot_roc_curve
19 )
20 from statsmodels.stats.outliers_influence import variance_inflation_factor
21 from imblearn.over_sampling import SMOTE
```



```
-----
ImportError                                Traceback (most recent call last)
<ipython-input-1-e16e9118419d> in <cell line: 15>()
    13 from sklearn.model_selection import train_test_split, KFold, cross_val_score
    14 from sklearn.preprocessing import MinMaxScaler
--> 15 from sklearn.metrics import (
    16     accuracy_score, confusion_matrix, classification_report,
    17     roc_auc_score, roc_curve, auc,

ImportError: cannot import name 'plot_confusion_matrix' from 'sklearn.metrics' (/usr/local/lib/python3.10/dist-packages/sklearn/metrics/__init__.py)
```

NOTE: If your import is failing due to a missing package, you can manually install dependencies using either !pip or !apt.

To view examples of installing some common dependencies, click the "Open Examples" button below.

OPEN EXAMPLES

Here is the information on this particular data set:

0. loan_amnt : The listed amount of the loan applied for by the borrower. If at some point in time, the credit department reduces the loan amount, then it will be reflected in this value.
1. term : The number of payments on the loan. Values are in months and can be either 36 or 60.
2. int_rate : Interest Rate on the loan
3. installment : The monthly payment owed by the borrower if the loan originates.
4. grade LC : assigned loan grade
5. sub_grade LC : assigned loan subgrade
6. emp_title : The job title supplied by the Borrower when applying for the loan.
7. emp_length : Employment length in years. Possible values are between 0 and 10 where 0 means less than one year and 10 means ten or more years.
8. home_ownership : The home ownership status provided by the borrower during registration or obtained from the credit report. Our values are: RENT, OWN, MORTGAGE, OTHER
9. annual_inc : The self-reported annual income provided by the borrower during registration.
10. verification_status : Indicates if income was verified by LC, not verified, or if the income source was verified
11. issue_d : The month which the loan was funded
12. loan_status : Current status of the loan
13. purpose : A category provided by the borrower for the loan request.
14. title : The loan title provided by the borrower
15. zip_code : The first 3 numbers of the zip code provided by the borrower in the loan application.
16. addr_state : The state provided by the borrower in the loan application

17. dti : A ratio calculated using the borrower's total monthly debt payments on the total debt obligations, excluding mortgage and the requested LC loan, divided by the borrower's self-reported monthly income.
18. earliest_cr_line : The month the borrower's earliest reported credit line was opened
19. open_acc : The number of open credit lines in the borrower's credit file.
20. pub_rec : Number of derogatory public records
21. revol_bal : Total credit revolving balance
22. revol_util : Revolving line utilization rate, or the amount of credit the borrower is using relative to all available revolving credit.
23. total_acc : The total number of credit lines currently in the borrower's credit file
24. initial_list_status : The initial listing status of the loan. Possible values are – W, F
25. application_type : Indicates whether the loan is an individual application or a joint application with two co-borrowers
26. mort_acc : Number of mortgage accounts.
27. pub_rec_bankruptcies : Number of public record bankruptcies

Reading the data file -

```
1 from google.colab import drive
2 drive.mount('/content/drive')
```

Mounted at /content/drive

```
1 !ls
```

```
drive  logistic_regression.csv  sample_data
```

```
1 data = pd.read_csv('logistic_regression.csv')
2 data.head()
```

annual_inc	...	open_acc	pub_rec	revol_bal	revol_util	total_acc	initial_list_status
117000.0	...	16.0	0.0	36369.0	41.8	25.0	w
65000.0	...	17.0	0.0	20131.0	53.3	27.0	f
43057.0	...	13.0	0.0	11987.0	92.2	26.0	f
54000.0	...	6.0	0.0	5472.0	21.5	13.0	f
55000.0	...	13.0	0.0	24584.0	69.8	43.0	f

```
1 # Shape of the dataset -
2 print("No. of rows: ", data.shape[0])
3 print("No. of columns: ", data.shape[1])
```

```
No. of rows: 231746
No. of columns: 27
```

```
1 # Checking the distribution of outcome labels -
2 data.loan_status.value_counts(normalize=True)*100
```

```
Fully Paid      80.423395
Charged Off     19.576605
Name: loan_status, dtype: float64
```

```
1 # Statistical summary of the dataset -
2 data.describe(include='all')
```

	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_title
count	231746.000000	231746	231746.000000	231746.000000	231746	231746	218320
unique	NaN	2	NaN	NaN	7	35	109777
top	NaN	36 months	NaN	NaN	B	B3	Teacher
freq	NaN	176621	NaN	NaN	67866	15546	2557
mean	14107.757955	NaN	13.640928	431.586377	NaN	NaN	NaN
std	8353.939311	NaN	4.466482	250.592582	NaN	NaN	NaN
min	500.000000	NaN	5.320000	16.250000	NaN	NaN	NaN
25%	8000.000000	NaN	10.490000	250.340000	NaN	NaN	NaN
50%	12000.000000	NaN	13.330000	375.430000	NaN	NaN	NaN
75%	20000.000000	NaN	16.490000	567.040000	NaN	NaN	NaN
max	40000.000000	NaN	30.990000	1533.810000	NaN	NaN	NaN

11 rows × 27 columns


```
1 data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 231746 entries, 0 to 231745
Data columns (total 27 columns):
#   Column                Non-Null Count  Dtype
---  -
0   loan_amnt             231746 non-null float64
1   term                  231746 non-null object
2   int_rate              231746 non-null float64
3   installment           231746 non-null float64
4   grade                 231746 non-null object
5   sub_grade             231746 non-null object
6   emp_title             218320 non-null object
7   emp_length            221005 non-null object
8   home_ownership        231746 non-null object
9   annual_inc            231746 non-null float64
10  verification_status   231746 non-null object
11  issue_d               231746 non-null object
12  loan_status           231746 non-null object
13  purpose               231746 non-null object
14  title                 230723 non-null object
15  dti                   231746 non-null float64
16  earliest_cr_line      231746 non-null object
17  open_acc              231746 non-null float64
18  pub_rec               231746 non-null float64
19  revol_bal             231746 non-null float64
20  revol_util            231576 non-null float64
21  total_acc             231746 non-null float64
22  initial_list_status   231746 non-null object
23  application_type      231746 non-null object
24  mort_acc              209558 non-null float64
25  pub_rec_bankruptcies  231426 non-null float64
26  address               231745 non-null object
dtypes: float64(12), object(15)
memory usage: 47.7+ MB
```

✦ Correlation Heatmap -

A correlation heatmap is a heatmap that shows a 2D correlation matrix between two discrete dimensions, using colored cells to represent data from usually a monochromatic scale. The values of the first dimension appear as the rows of the table while of the second dimension as a column. The color of the cell is proportional to the number of measurements that match the dimensional value. This makes correlation heatmaps ideal for data analysis since it makes patterns easily readable and highlights the differences and variation in the same data. A correlation heatmap, like a regular heatmap, is assisted by a colorbar making data easily readable and comprehensible.

```
1 plt.figure(figsize=(12, 8))
2 sns.heatmap(data.corr(method='spearman'), annot=True, cmap='viridis')
3 plt.show()
```




We noticed almost perfect correlation between "loan_amnt" the "installment" feature.

- installment: The monthly payment owed by the borrower if the loan originates.
- loan_amnt: The listed amount of the loan applied for by the borrower. If at some point in time, the credit department reduces the loan amount, then it will be reflected in this value.

So, we can drop either one of those columns.

```
1 data.drop(columns=['installment'], axis=1, inplace=True)

1 plt.figure(figsize=(12, 8))
2 sns.heatmap(data.corr(method='spearman'), annot=True, cmap='viridis')
3 plt.show()
```



✓ Data Exploration -

1. The no of people those who have fully paid are 318357 and that of Charged Off are 77673.

```
1 data.groupby(by='loan_status')['loan_amnt'].describe()
```

	count	mean	std	min	25%	50%	75%	max
loan_status								
Charged Off	45368.0	15135.654536	8484.848789	1000.0	8700.0	14000.0	20000.0	40000.0
Fully Paid	186378.0	13857.548101	8302.548155	500.0	7500.0	12000.0	19200.0	40000.0

2. The majority of people have home ownership as Mortgage and Rent.

```
1 data['home_ownership'].value_counts()
```

```
MORTGAGE    116104
RENT         93551
OWN          22010
OTHER         59
NONE         19
ANY           3
Name: home_ownership, dtype: int64
```

3. Combining the minority classes as 'OTHER'.

```
1 data.loc[(data.home_ownership == 'ANY') | (data.home_ownership == 'NONE'), 'home_ownership'] = 'OTHER'
2 data.home_ownership.value_counts()
```

```
MORTGAGE    116104
RENT         93551
OWN          22010
OTHER         81
Name: home_ownership, dtype: int64
```

```
1 data['home_ownership'].value_counts()
```

```

MORTGAGE    116104
RENT        93551
OWN         22010
OTHER       81
Name: home_ownership, dtype: int64

```

```

1 # Checking the distribution of 'Other' -
2 data.loc[data['home_ownership']=='OTHER', 'loan_status'].value_counts()

```

```

Fully Paid    64
Charged Off   17
Name: loan_status, dtype: int64

```

4. Covertng string to date-time format.

```

1 data['issue_d'] = pd.to_datetime(data['issue_d'])
2 data['earliest_cr_line'] = pd.to_datetime(data['earliest_cr_line'])

```

5. Saw some issues in title (Looks like it was filled manually and needs some fixing).

```

1 data['title'].value_counts()[:20]

```

```

Debt consolidation      89183
Credit card refinancing 30086
Home improvement        8932
Other                   7597
Debt Consolidation      6854
Major purchase          2887
Consolidation           2305
debt consolidation      2055
Business                1701
Debt Consolidation Loan 1673
Medical expenses        1605
Car financing           1242
Credit Card Consolidation 1049
Vacation                1025
Moving and relocation    986
consolidation           971
Personal Loan           913
Home Improvement        746
Consolidation Loan      739
Home buying             674
Name: title, dtype: int64

```

```

1 data['title'] = data.title.str.lower()

```

```

1 data.title.value_counts()[:10]

```

```

debt consolidation      98370
credit card refinancing 30256
home improvement        10011
other                   7637
consolidation           3344
major purchase          3024
debt consolidation loan 2059
business                1749
medical expenses        1652
credit card consolidation 1567
Name: title, dtype: int64

```

✓ Visualization -


The grade of majority of people those who have fully paid the loan is 'B' and have subgrade 'B3'.

So from where we can infer that people with grade 'B' and subgrade 'B3' are more likely to fully pay the loan.

```

1 plt.figure(figsize=(15, 10))
2
3 plt.subplot(2, 2, 1)
4 grade = sorted(data.grade.unique().tolist())
5 sns.countplot(x='grade', data=data, hue='loan_status', order=grade)
6
7 plt.subplot(2, 2, 2)
8 sub_grade = sorted(data.sub_grade.unique().tolist())
9 g = sns.countplot(x='sub_grade', data=data, hue='loan_status', order=sub_grade)
10 g.set_xticklabels(g.get_xticklabels(), rotation=90);

```




1 grade

```
['A', 'B', 'C', 'D', 'E', 'F', 'G']
```

```

1 plt.figure(figsize=(15, 20))
2
3 plt.subplot(4, 2, 1)
4 sns.countplot(x='term', data=data, hue='loan_status')
5
6 plt.subplot(4, 2, 2)
7 sns.countplot(x='home_ownership', data=data, hue='loan_status')
8
9 plt.subplot(4, 2, 3)
10 sns.countplot(x='verification_status', data=data, hue='loan_status')
11
12 plt.subplot(4, 2, 4)
13 g = sns.countplot(x='purpose', data=data, hue='loan_status')
14 g.set_xticklabels(g.get_xticklabels(), rotation=90);

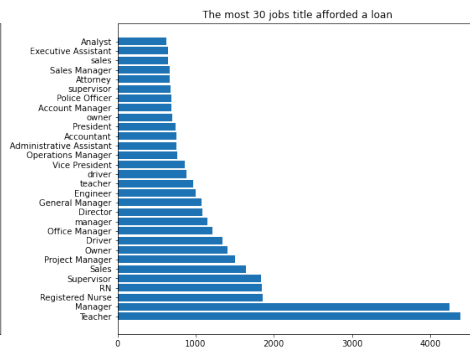

```



1 Start coding or [generate](#) with AI.

Manager and Teacher are the most afforded loan job titles.

```
1 plt.figure(figsize=(15, 12))
2
3 plt.subplot(2, 2, 1)
4 order = ['< 1 year', '1 year', '2 years', '3 years', '4 years', '5 years',
5         '6 years', '7 years', '8 years', '9 years', '10+ years',]
6 g = sns.countplot(x='emp_length', data=data, hue='loan_status', order=order)
7 g.set_xticklabels(g.get_xticklabels(), rotation=90);
8
9 plt.subplot(2, 2, 2)
10 plt.barh(data.emp_title.value_counts()[:30].index, data.emp_title.value_counts()[:30])
11 plt.title("The most 30 jobs title afforded a loan")
12 plt.tight_layout()
```

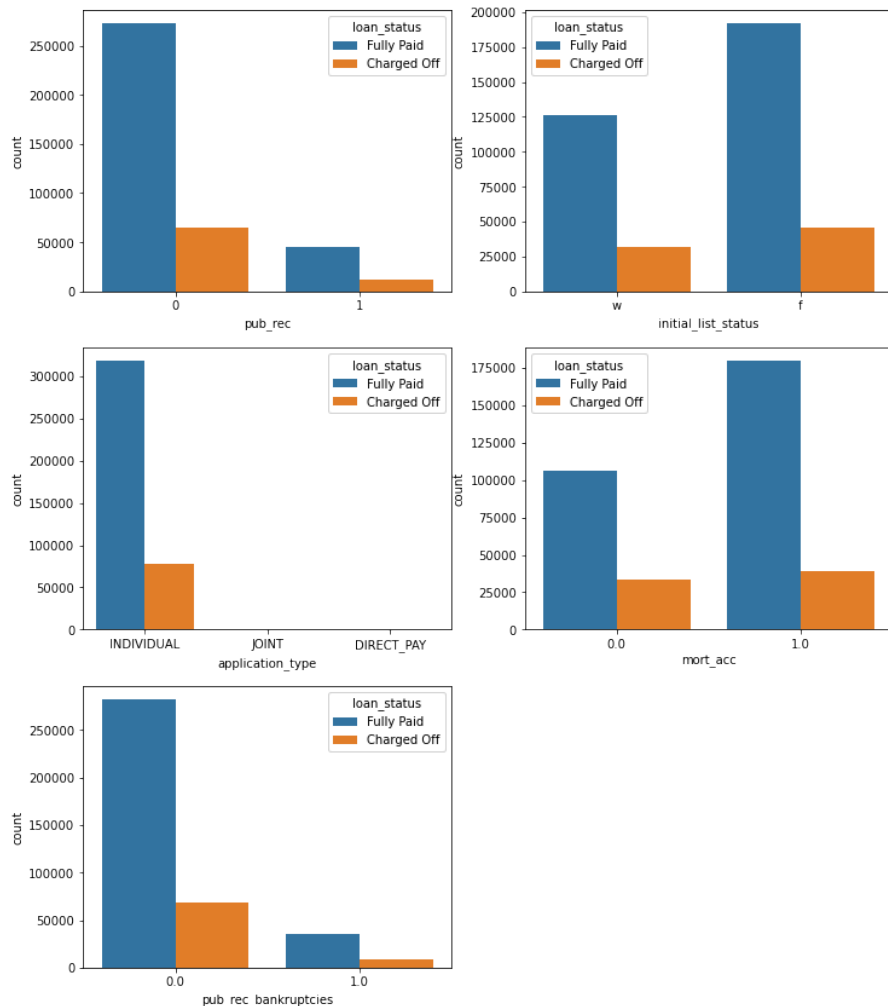



Feature Engineering -

```

1 def pub_rec(number):
2     if number == 0.0:
3         return 0
4     else:
5         return 1
6
7 def mort_acc(number):
8     if number == 0.0:
9         return 0
10    else:
11        return 1
12
13
14 def pub_rec_bankruptcies(number):
15     if number == 0.0:
16         return 0
17     else:
18         return 1
19
20
21 data['pub_rec'] = data.pub_rec.apply(pub_rec)
22 data['mort_acc'] = data.mort_acc.apply(mort_acc)
23 data['pub_rec_bankruptcies'] = data.pub_rec_bankruptcies.apply(pub_rec_bankruptcies)
24
25
26 plt.figure(figsize=(12, 30))
27
28 plt.subplot(6, 2, 1)
29 sns.countplot(x='pub_rec', data=data, hue='loan_status')
30
31 plt.subplot(6, 2, 2)
32 sns.countplot(x='initial_list_status', data=data, hue='loan_status')
33
34 plt.subplot(6, 2, 3)
35 sns.countplot(x='application_type', data=data, hue='loan_status')
36
37 plt.subplot(6, 2, 4)
38 sns.countplot(x='mort_acc', data=data, hue='loan_status')
39
40 plt.subplot(6, 2, 5)
41 sns.countplot(x='pub_rec_bankruptcies', data=data, hue='loan_status')
42
43 plt.show()

```



```
1 # Mapping of target variable -
2 data['loan_status'] = data.loan_status.map({'Fully Paid':0, 'Charged Off':1})
```

```
1 data.isnull().sum()/len(data)*100
```

```

loan_amnt      0.000000
term           0.000000
int_rate       0.000000
grade          0.000000
sub_grade      0.000000
emp_title      5.789208
emp_length     4.621115
home_ownership 0.000000
annual_inc     0.000000
verification_status 0.000000
issue_d        0.000000
loan_status    0.000000
purpose        0.000000
title          0.443148
dti            0.000000
earliest_cr_line 0.000000
open_acc       0.000000
pub_rec        0.000000
revol_bal      0.000000
revol_util     0.069692
total_acc      0.000000
initial_list_status 0.000000
application_type 0.000000
mort_acc       9.543469
pub_rec_bankruptcies 0.135091
address        0.000000
dtype: float64
```

Very Important: Mean Target Imputation

```
1 data.groupby(by='total_acc').mean()
```

	loan_amnt	int_rate	annual_inc	loan_status	dti	open_acc	pub_rec
total_acc							
2.0	6672.222222	15.801111	64277.777778	0.222222	2.279444	1.611111	0.001
3.0	6042.966361	15.615566	41270.753884	0.220183	6.502813	2.611621	0.031
4.0	7587.399031	15.069491	42426.565969	0.214055	8.411963	3.324717	0.031
5.0	7845.734714	14.917564	44394.098003	0.203156	10.118328	3.921598	0.051
6.0	8529.019843	14.651752	48470.001156	0.215874	11.222542	4.511119	0.071
...
124.0	23200.000000	17.860000	66000.000000	1.000000	14.040000	43.000000	0.001
129.0	25000.000000	7.890000	200000.000000	0.000000	8.900000	48.000000	0.001
135.0	24000.000000	15.410000	82000.000000	0.000000	33.850000	57.000000	0.001
150.0	35000.000000	8.670000	189000.000000	0.000000	6.630000	40.000000	0.001
151.0	35000.000000	13.990000	160000.000000	1.000000	12.650000	26.000000	0.001

118 rows × 11 columns

```
1 total_acc_avg = data.groupby(by='total_acc').mean().mort_acc
2 # Saving mean of mort_acc according to total_acc_avg (you can pick any variable for your understanding)
```

```
1 def fill_mort_acc(total_acc, mort_acc):
2     if np.isnan(mort_acc):
3         return total_acc_avg[total_acc].round()
4     else:
5         return mort_acc
```

```
1 data['mort_acc'] = data.apply(lambda x: fill_mort_acc(x['total_acc'], x['mort_acc']), axis=1)
```

```
1 data.isnull().sum()/len(data)*100
```

```
loan_amnt      0.000000
term           0.000000
int_rate       0.000000
grade          0.000000
sub_grade      0.000000
emp_title      5.789208
emp_length     4.621115
home_ownership 0.000000
annual_inc     0.000000
verification_status 0.000000
issue_d        0.000000
loan_status    0.000000
purpose        0.000000
title          0.443148
dti            0.000000
earliest_cr_line 0.000000
open_acc       0.000000
pub_rec        0.000000
revol_bal      0.000000
revol_util     0.069692
total_acc      0.000000
initial_list_status 0.000000
application_type 0.000000
mort_acc       0.000000
pub_rec_bankruptcies 0.135091
address        0.000000
dtype: float64
```

```
1 # Current no. of rows -
```

```
2 data.shape
```

```
(396030, 26)
```

```
1 # Dropping rows with null values -  
2 data.dropna(inplace=True)
```

```
1 # Remaining no. of rows -  
2 data.shape
```


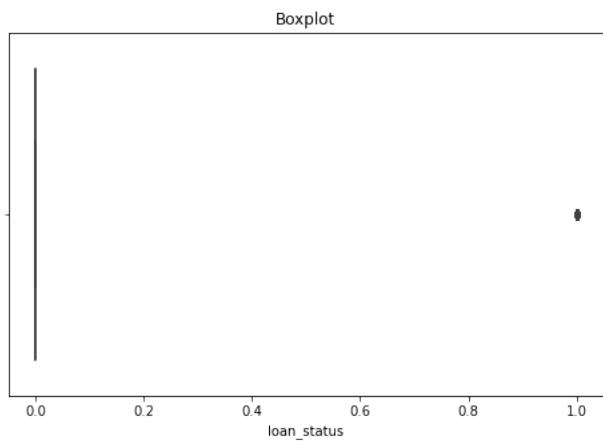
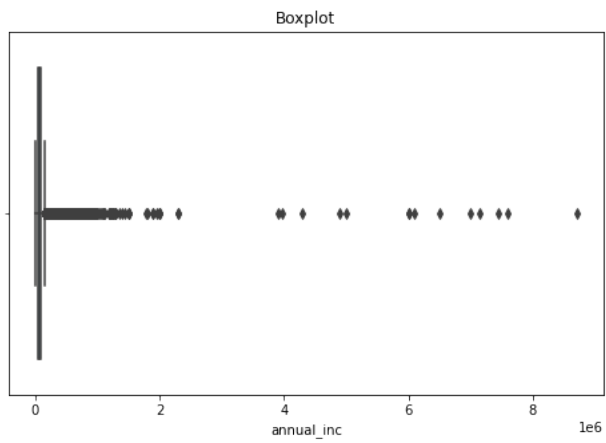


```
(370622, 26)
```


✓ Outlier Detection & Treatment -

```
1 numerical_data = data.select_dtypes(include='number')  
2 num_cols = numerical_data.columns  
3 len(num_cols)
```


```
12
```

```
1 def box_plot(col):  
2     plt.figure(figsize=(8, 5))  
3     sns.boxplot(x=data[col])  
4     plt.title('Boxplot')  
5     plt.show()  
6  
7 for col in num_cols:  
8     box_plot(col)
```







Boxplot




Boxplot

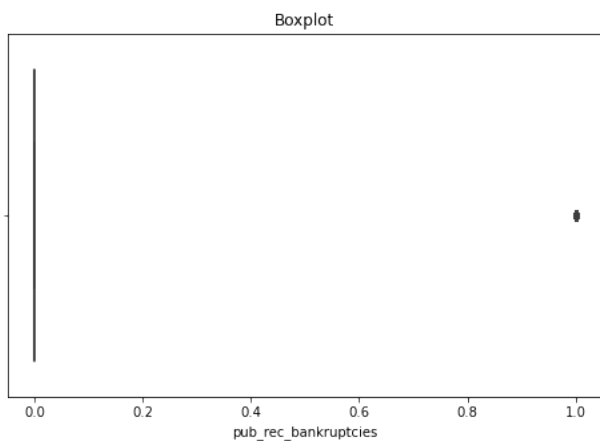
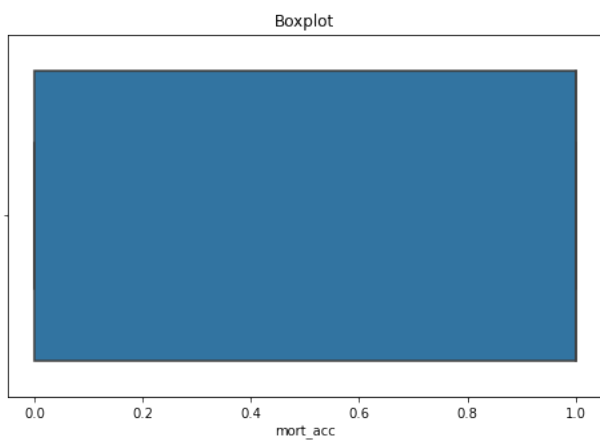




Boxplot



Boxplot





```

1 for col in num_cols:
2     mean = data[col].mean()
3     std = data[col].std()
4
5     upper_limit = mean+3*std
6     lower_limit = mean-3*std
7
8     data = data[(data[col]<upper_limit) & (data[col]>lower_limit)]
9
10 data.shape

(354519, 26)

```

✓ Data Preprocessing -

```

1 # Term -
2 data.term.unique()

```

```
array([' 36 months', ' 60 months'], dtype=object)
```

```
1 Start coding or generate with AI.
```

```
1 term_values = {' 36 months': 36, ' 60 months': 60}
2 data['term'] = data.term.map(term_values)
```

```
1 # Initial List Status -
2 data['initial_list_status'].unique()
```

```
array(['w', 'f'], dtype=object)
```

```
1 list_status = {'w': 0, 'f': 1}
2 data['initial_list_status'] = data.initial_list_status.map(list_status)
```

```
1 # Let's fetch ZIP from address and then drop the remaining details -
2 data['zip_code'] = data.address.apply(lambda x: x[-5:])
```

```
1 data['zip_code'].value_counts(normalize=True)*100
```

```
70466    14.382022
30723    14.277373
22690    14.268347
48052    14.127028
00813    11.610097
29597    11.537322
05113    11.516731
93700     2.774746
11650     2.772771
86630     2.733563
Name: zip_code, dtype: float64
```

```
1 # Dropping some variables which IMO we can let go for now -
2 data.drop(columns=['issue_d', 'emp_title', 'title', 'sub_grade',
3                   'address', 'earliest_cr_line', 'emp_length'],
4           axis=1, inplace=True)
```

✓ One-hot Encoding -

```
1 dummies = ['purpose', 'zip_code', 'grade', 'verification_status', 'application_type', 'home_ownership']
2 data = pd.get_dummies(data, columns=dummies, drop_first=True)
```

```
1 pd.set_option('display.max_columns', None)
2 pd.set_option('display.max_rows', None)
3
4 data.head()
```

	loan_amnt	term	int_rate	annual_inc	loan_status	dti	open_acc	pub_rec	revol_ba
0	10000.0	36	11.44	117000.0	0	26.24	16.0	0	36369
1	8000.0	36	11.99	65000.0	0	22.05	17.0	0	20131
2	15600.0	36	10.49	43057.0	0	12.79	13.0	0	11987
3	7200.0	36	6.49	54000.0	0	2.60	6.0	0	5472
4	24375.0	60	17.27	55000.0	1	33.95	13.0	0	24584

```
1 data.shape
(354519, 49)
```

✓ Data Preparation for Modeling -


```

1 X = data.drop('loan_status', axis=1)
2 y = data['loan_status']

1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30,
2                                                    stratify=y, random_state=42)

1 print(X_train.shape)
2 print(X_test.shape)

(248163, 48)
(106356, 48)

```

✓ MinMaxScaler -

For each value in a feature, MinMaxScaler subtracts the minimum value in the feature and then divides by the range. The range is the difference between the original maximum and original minimum.

MinMaxScaler preserves the shape of the original distribution. It doesn't meaningfully change the information embedded in the original data.

```

1 scaler = MinMaxScaler()
2 X_train = scaler.fit_transform(X_train)
3 X_test = scaler.transform(X_test)

```

✓ Logistic Regression

```

1 logreg = LogisticRegression(max_iter=1000)
2 logreg.fit(X_train, y_train)

LogisticRegression(max_iter=1000)

1 y_pred = logreg.predict(X_test)
2 print('Accuracy of Logistic Regression Classifier on test set: {:.3f}'.format(logreg.score(X_test, y_test)))

Accuracy of Logistic Regression Classifier on test set: 0.890

```

✓ Confusion Matrix -

```

1 confusion_matrix = confusion_matrix(y_test, y_pred)
2 print(confusion_matrix)

[[85365  523]
 [11131  9337]]

```

✓ Classification Report -

```

1 print(classification_report(y_test, y_pred))

              precision    recall  f1-score   support

     0               0.88       0.99       0.94       85888
     1               0.95       0.46       0.62       20468

 accuracy               0.89       106356
 macro avg              0.92       0.73       0.78       106356
 weighted avg           0.90       0.89       0.87       106356

```

ROC Curve -

An ROC curve (receiver operating characteristic curve) is a graph showing the performance of a classification model at all classification thresholds. This curve plots two parameters:

- True Positive Rate
- False Positive Rate

True Positive Rate (TPR) is a synonym for recall and is therefore defined as follows:

- $TPR = (TP) / (TP + FN)$

False Positive Rate (FPR) is defined as follows:

- $FPR = (FP) / (FP + TN)$


An ROC curve plots TPR vs. FPR at different classification thresholds. Lowering the classification threshold classifies more items as positive, thus increasing both False Positives and True Positives. The following figure shows a typical ROC curve.

✓ AUC (Area under the ROC Curve) -


AUC stands for "Area under the ROC Curve." That is, AUC measures the entire two-dimensional area underneath the entire ROC curve (think integral calculus) from (0,0) to (1,1).

AUC provides an aggregate measure of performance across all possible classification thresholds. One way of interpreting AUC is as the probability that the model ranks a random positive example more highly than a random negative example. For example, given the following examples, which are arranged from left to right in ascending order of logistic regression predictions:

```
1 logit_roc_auc = roc_auc_score(y_test, logreg.predict(X_test))
2 fpr, tpr, thresholds = roc_curve(y_test, logreg.predict_proba(X_test)[: ,1])
3 plt.figure()
4 plt.plot(fpr, tpr, label='Logistic Regression (area = %0.2f)' % logit_roc_auc)
5 plt.plot([0, 1], [0, 1], 'r--')
6 plt.xlim([0.0, 1.0])
7 plt.ylim([0.0, 1.05])
8 plt.xlabel('False Positive Rate')
9 plt.ylabel('True Positive Rate')
10 plt.title('Receiver operating characteristic')
11 plt.legend(loc="lower right")
12 plt.savefig('Log_ROC')
13 plt.show()
```



```
1 def precision_recall_curve_plot(y_test, pred_proba_c1):
2     precisions, recalls, thresholds = precision_recall_curve(y_test, pred_proba_c1)
3
4     threshold_boundary = thresholds.shape[0]
5     # plot precision
6     plt.plot(thresholds, precisions[0:threshold_boundary], linestyle='--', label='precision')
7     # plot recall
8     plt.plot(thresholds, recalls[0:threshold_boundary], label='recalls')
9
10    start, end = plt.xlim()
11    plt.xticks(np.round(np.arange(start, end, 0.1), 2))
12
13    plt.xlabel('Threshold Value'); plt.ylabel('Precision and Recall Value')
14    plt.legend(); plt.grid()
15    plt.show()
16
17 precision_recall_curve_plot(y_test, logreg.predict_proba(X_test)[: ,1])
```



✓ Multicollinearity check using Variance Inflation Factor (VIF) -

Multicollinearity occurs when two or more independent variables are highly correlated with one another in a regression model. Multicollinearity can be a problem in a regression model because we would not be able to distinguish between the individual effects of the independent variables on the dependent variable.

Multicollinearity can be detected via various methods. One such method is Variance Inflation Factor aka VIF. In VIF method, we pick each independent feature and regress it against all of the other independent features. VIF score of an independent variable represents how well the variable is explained by other independent variables.

$VIF = 1/(1-R^2)$

```
1 def calc_vif(X):
2     # Calculating the VIF
3     vif = pd.DataFrame()
4     vif['Feature'] = X.columns
5     vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
6     vif['VIF'] = round(vif['VIF'], 2)
7     vif = vif.sort_values(by='VIF', ascending = False)
8     return vif
9
10 calc_vif(X)[:5]
```

	Feature	VIF
43	application_type_INDIVIDUAL	156.97
2	int_rate	122.82
14	purpose_debt_consolidation	51.00
1	term	27.30
13	purpose_credit_card	18.48

```
1 X.drop(columns=['application_type_INDIVIDUAL'], axis=1, inplace=True)
2 calc_vif(X)[:5]
```

	Feature	VIF
2	int_rate	103.43
14	purpose_debt_consolidation	27.49
1	term	24.31
5	open_acc	13.75
9	total_acc	12.69

```
1 X.drop(columns=['int_rate'], axis=1, inplace=True)
2 calc_vif(X)[:5]
```

	Feature	VIF
1	term	23.35
13	purpose_debt_consolidation	22.35
4	open_acc	13.64
8	total_acc	12.69
7	revol_util	9.06